

COMP 5070

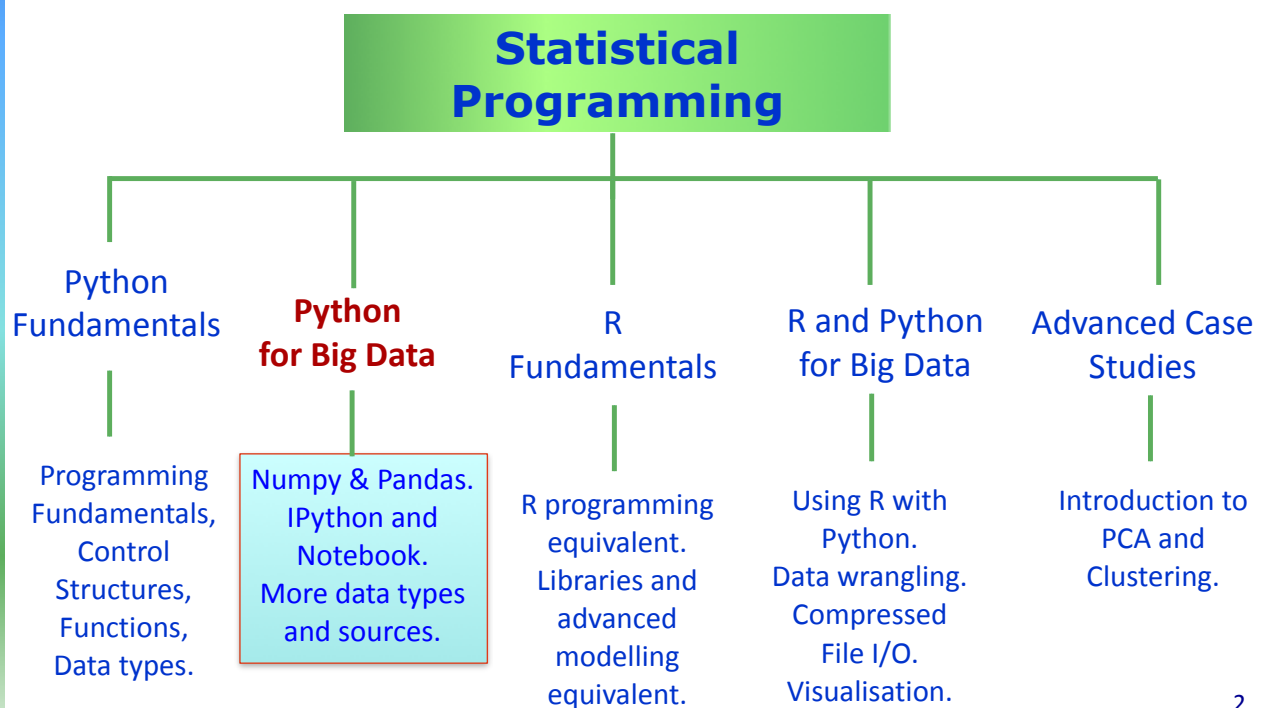
**Statistical Programming
for
Data Science**

Python for Big Data:
pandas

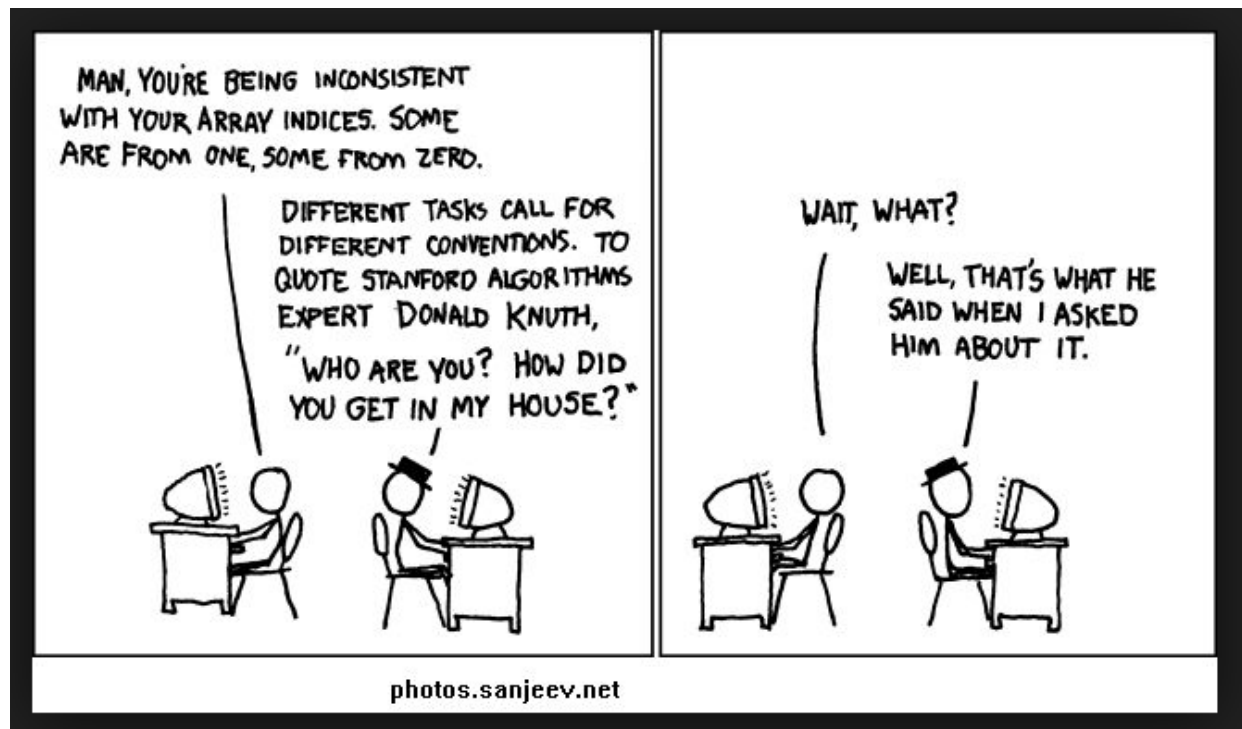


University of
South Australia

The course at a glance



The issue with pandas ...



First (and only) up ...

pandas

- The answer to R?
- New (and rich) data structures:
 - Series
 - DataFrame
- Indexing:
 - Regular
 - Hierarchical, Partial
- Immutability and Reindexing
- Function mapping
 - apply, applymap
- Missing data



pandas in a Nutshell

- "R for Python"
- Provides easy to use data structures and many useful helper functions for data cleanup and transformations.
- Fast! (backed by numpy arrays)
- Contains high-level data structures and manipulation tools including structured or tabular data.
- Provides a rich, high-level interface making most common data tasks very concise and simple.
- Provides domain-specific functionality, e.g. time series manipulation and easy handling of missing data (not present in NumPy).

5

pandas in a Nutshell

- A clean axis indexing design to support fast data alignment, lookups, hierarchical indexing, and more high-performance data structures.
- Functionality includes:
 - Series/TimeSeries: 1D labelled vector
 - DataFrame: 2D spreadsheet-like structure
 - Panel: 3D labeled array, collection of DataFrames
- SQL-like functionality: GroupBy, joining/merging, etc. Missing data handling.
- Etymology: panel data structures

6

Indexing: The pandas Killer Feature

- Each axis has an index.
- **Automatic alignment between differently-indexed objects:** makes it nearly impossible to accidentally combine misaligned data
- **Hierarchical indexing** provides an intuitive way of structuring and working with higher-dimensional data
- Natural way of expressing “group by” and join-type operations
- **Claim:** considered to be a better integrated and more flexible indexing than anything available in R or MATLAB

7

Pandas Data Structures

- The development of pandas introduces two new data structures to Python:
 1. **Series**
 2. **DataFrame**
- Both are built on top of NumPy (so it's fast!).
- Usually import pandas with the shorthand **pd**, i.e.

```
import pandas as pd
```

8

The Series Data Structure

- A **Series** is a one-dimensional object similar to an array, list, or column in a table. It is a mutable structure.
- Series allows for a labelled index to be assigned to each item.
- As with the List structure, by default, **each item in a Series will receive an index label starting from 0.**
- The simplest Series is formed from only an array of data:

```
In [3]: obj = pd.Series([4, 7, -5, 3])
Out[3]:
0      4
1      7
2     -5
3      3
dtype: int64
```

Index is shown on the left,
values on the right!

9

Series – extracting indices and value

- Can **extract the indices and values** using `obj.index` and `obj.values` respectively, i.e.

```
In [4]: obj.values
Out[4]: array([ 4,  7, -5,  3])
```

```
In [5]: obj.index
Out[5]: Int64Index([0, 1, 2, 3], dtype='int64')
```

- Can also **specify the desired indices:**

```
In [8]: obj2=pd.Series([4, 7, -5, 3],index=['d', 'b', 'a', 'c'])
Out[8]:
d      4
b      7
a     -5
c      3
dtype: int64
```

10

Series – selecting values

- Can use values in the index to select values:

```
In [9]: obj2['a']
```

```
Out[9]:
```

```
-5
```

```
In [11]:
```

```
obj2['d'] = 6
```

```
obj2[['c', 'a', 'd']]
```

```
Out[11]:
```

```
c 3
```

```
a -5
```

```
d 6
```

```
dtype: int64
```

```
In [12]: obj2[obj2 > 0]
```

```
Out[12]:
```

```
d 6
```

```
b 7
```

```
c 3
```

```
dtype: int64
```

```
In [14]:
```

```
obj2*2
```

```
Out[14]:
```

```
d 12 b 14 a -10 c 6 dtype:
```

```
int64
```

- Note that using e.g. NumPy boolean indexing, scalar multiplication or ufuncs (aggregation) will preserve the indexing.

11

Converting dict{} to Series

- Can think about a Series as a fixed-length, ordered dict, as it is a mapping of index values to data values:

```
In [15]: 'b' in obj2
```

```
Out[15]:
```

```
True
```

```
In [16]: 'e' in obj2
```

```
Out[16]:
```

```
False
```

```
In [17]: sdata = {'Ohio': 35000,
```

```
'Texas': 71000, 'Oregon': 16000,
```

```
'Utah': 5000}
```

```
In [19]: obj3 = pd.Series(sdata)
```

```
obj3
```

```
Out[19]:
```

```
Ohio    35000
```

```
Oregon  16000
```

```
Texas   71000
```

```
Utah     5000
```

- Can create a Series from a dictionary by calling the `panda.Series()` function. 12

Differently Indexed Data

- Series automatically aligns differently-indexed data in arithmetic operations:

```
In [25]: obj3
Out[25]:
```

```
Ohio 35000
Oregon 16000
Texas 71000
Utah 5000
```

```
In [26]: obj4
Out[26]:
```

```
California NaN
Ohio 35000
Oregon 16000
Texas 71000
```

```
In [27]: obj3+obj4
Out[27]:
```

```
California      NaN
Ohio            70000
Oregon          32000
Texas          142000
Utah            NaN
```

13

DataFrame

- The DataFrame is designed to be similar to the R dataframe structure.
- Represents a **tabular, spreadsheet-like data structure** containing an **ordered collection of columns**, each of which can be a **different value type** (numeric, string, boolean, etc.)
- The **DataFrame has both a row and column index**.
- While a DataFrame stores the data internally in a two-dimensional format, **can represent much higher-dimensional data in a tabular format using hierarchical indexing**.
- Hierarchical indexing (more later in the course!) is a key component in many of the more advanced data-handling features in pandas.

14

Data inputs to DataFrame

| Type | Notes |
|----------------------------------|---|
| 2D ndarray | A matrix of data, passing optional row and column labels |
| dict of arrays, lists, or tuples | Each sequence becomes a column in the DataFrame. All sequences must be the same length. |
| NumPy structured/record array | Treated as the "dict of arrays" case |
| dict of Series | Each value becomes a column. Indexes from each Series are unioned together to form the result's row index if no explicit index is passed. |
| dict of dicts | Each inner dict becomes a column. Keys are unioned to form the row index as in the "dict of Series" case. |
| list of dicts or Series | Each item becomes a row in the DataFrame. Union of dict keys or Series indexes become the DataFrame's column labels |
| List of lists or tuples | Treated as the "2D ndarray" case |
| Another DataFrame | The DataFrame's indexes are used unless different ones are passed |
| NumPy MaskedArray | Like the "2D ndarray" case except masked values become NA/missing in the DataFrame result |

15

DataFrame - Construction

- Many ways to construct a DataFrame. A common way is to convert a dict structure using DataFrame():

```
In [30]: data = {'state': ['Ohio', 'Ohio', 'Ohio', 'Nevada',  
                          'Nevada'], 'year': [2000, 2001, 2002, 2001, 2002],  
              'pop': [1.5, 1.7, 3.6, 2.4, 2.9]}  
frame = pd.DataFrame(data)
```

Out[30]:

| | pop | state | year |
|---|-----|--------|------|
| 0 | 1.5 | Ohio | 2000 |
| 1 | 1.7 | Ohio | 2001 |
| 2 | 3.6 | Ohio | 2002 |
| 3 | 2.4 | Nevada | 2001 |
| 4 | 2.9 | Nevada | 2002 |

5 rows x 3 columns

16

DataFrame – Adding columns

- Can specify a sequence of columns and if the column doesn't exist, you will obtain NAs:

```
In [34]:  
frame2 = pd.DataFrame(data, columns=['year', 'state', 'pop', 'debt'],  
index=['one', 'two', 'three', 'four', 'five'])
```

Out[34]:

| | year | state | pop | debt |
|-------|------|--------|-----|------|
| one | 2000 | Ohio | 1.5 | NaN |
| two | 2001 | Ohio | 1.7 | NaN |
| three | 2002 | Ohio | 3.6 | NaN |
| four | 2001 | Nevada | 2.4 | NaN |
| five | 2002 | Nevada | 2.9 | NaN |

17

DataFrame – Data Retrieval

- Can retrieve a column by name

```
In [35]: frame2['state']  
Out[35]:  
one Ohio  
two Ohio  
three Ohio  
four Nevada  
five Nevada  
Name: state, dtype: object
```

*The column returned
when indexing a
DataFrame is a view on
the underlying data, not
a copy.*

or attribute (use . notation, very R-like!):

```
In [36]: frame2.year  
Out[36]:  
one 2000  
two 2001  
three 2002  
four 2001  
five 2002  
Name: year, dtype: int64
```

*To create a copy, use
pd.Series.copy()*

18

DataFrame – Data Retrieval

- Rows can be retrieved by position, name or other methods, e.g. `ix`

```
In [37]: frame2.ix['three']
```

```
Out[37]:  
year 2002  
state Ohio  
pop 3.6  
debt NaN  
Name: three, dtype: object
```

- Columns can be easily reassigned values:

```
In [40]: frame2['debt'] = np.arange(5.)
```

| | year | state | pop | debt |
|-------|------|--------|-----|------|
| one | 2000 | Ohio | 1.5 | 0 |
| two | 2001 | Ohio | 1.7 | 1 |
| three | 2002 | Ohio | 3.6 | 2 |
| four | 2001 | Nevada | 2.4 | 3 |
| five | 2002 | Nevada | 2.9 | 4 |

DataFrame – Inserting Data

- If assigning a list or array to a column in a dataframe, if the sizes do not match exactly, NAs will be used to fill in the gaps:

```
In [41]: val = pd.Series([-1.2, -1.5, -1.7],  
                        index=['two', 'four', 'five'])  
frame2['debt'] = val
```

Out[42]:

| | year | state | pop | debt |
|-------|------|--------|-----|------|
| one | 2000 | Ohio | 1.5 | NaN |
| two | 2001 | Ohio | 1.7 | -1.2 |
| three | 2002 | Ohio | 3.6 | NaN |
| four | 2001 | Nevada | 2.4 | -1.5 |
| five | 2002 | Nevada | 2.9 | -1.7 |

DataFrame – Deleting Columns

- Assigning a column that doesn't exist will create a new column. The `del` keyword will delete columns as with a dict.

```
In [43]: frame2['eastern'] = frame2.state == 'Ohio'
```

Out[43]:

| | year | state | pop | debt | eastern |
|-------|------|--------|-----|------|---------|
| one | 2000 | Ohio | 1.5 | NaN | True |
| two | 2001 | Ohio | 1.7 | -1.2 | True |
| three | 2002 | Ohio | 3.6 | NaN | True |
| four | 2001 | Nevada | 2.4 | -1.5 | False |
| five | 2002 | Nevada | 2.9 | -1.7 | False |

```
In [44]: del frame2['eastern']  
         frame2.columns
```

Out[44]:

```
Index(['year', 'state', 'pop', 'debt'], dtype='object')
```

21

Dropping entries from an axis

- Dropping one or more entries from an axis is relatively straightforward.
Can use a **Series** or **DataFrame** approach:

```
In [83]:
```

```
obj = pd.Series(np.arange(5.), index=['a', 'b', 'c', 'd', 'e'])  
new_obj = obj.drop('c')
```

```
In [85]: obj
```

Out[85]:

```
a 0  
b 1  
c 2  
d 3  
e 4
```

```
In [86]: new_obj
```

Out[86]:

```
a 0  
b 1  
d 3  
e 4
```

22

Dropping entries from an axis

```
data =  
pd.DataFrame(np.arange(16).reshape((4, 4)),  
index=['Ohio', 'Colorado', 'Utah', 'New York'],  
columns=['one', 'two', 'three', 'four'])
```

| | one | two | three | four |
|----------|-----|-----|-------|------|
| Ohio | 0 | 1 | 2 | 3 |
| Colorado | 4 | 5 | 6 | 7 |
| Utah | 8 | 9 | 10 | 11 |
| New York | 12 | 13 | 14 | 15 |

```
data.drop(['Colorado', 'Ohio'])
```

| | one | two | three | four |
|----------|-----|-----|-------|------|
| Utah | 8 | 9 | 10 | 11 |
| New York | 12 | 13 | 14 | 15 |

```
data.drop(['two', 'four'], axis=1)
```

| | one | three |
|----------|-----|-------|
| Ohio | 0 | 2 |
| Colorado | 4 | 6 |
| Utah | 8 | 10 |
| New York | 12 | 14 |

23

DataFrame – Creating Nested Dicts{}

- Can also create nested dicts:

```
pop = {'Nevada': {2001: 2.4, 2002: 2.9},  
       'Ohio': {2000: 1.5, 2001: 1.7, 2002: 3.6}}  
frame3 = pd.DataFrame(pop)
```

Out[47]:

| | Nevada | Ohio |
|------|--------|------|
| 2000 | NaN | 1.5 |
| 2001 | 2.4 | 1.7 |
| 2002 | 2.9 | 3.6 |

*Outer keys will become
columns. Inner keys will
become rows.*

```
frame3.T
```

Out[48]:

| | 2000 | 2001 | 2002 |
|--------|------|------|------|
| Nevada | NaN | 2.4 | 2.9 |
| Ohio | 1.5 | 1.7 | 3.6 |

*DataFrames can be
transposed!*

24

DataFrames – Setting Row/Column Names

- If a DataFrame's index and columns have their name attributes set, these will also be displayed:

```
In [49]:  
frame3.index.name = 'year'; frame3.columns.name = 'state'  
frame3
```

Out[49]:

| state | Nevada | Ohio |
|-------|--------|------|
| year | | |
| 2000 | NaN | 1.5 |
| 2001 | 2.4 | 1.7 |
| 2002 | 2.9 | 3.6 |

- Like Series, can use a `values` attribute (e.g. `frame3.values`) and if the column types are all different, dtype will be: `dtype=object`.

25

Merging Data Frames (join)

- `pandas.merge` allows two DataFrames to be joined on one or more keys.
- `pandas.merge` operates as an inner join.
- Can change this using the option **how**
- **how** allows two data frames to be joined with the options left, right, outer and inner, which tells pandas:
 - **left**: use only keys from left frame (SQL: left outer join)
 - **right**: use only keys from right frame (SQL: right outer join)
 - **outer**: use union of keys from both frames (SQL: full outer join)
 - **inner**: use intersection of keys from both frames (SQL: inner join)

26

Merging Data Frames (join)

```
left_frame = pd.DataFrame({'key': range(5), 'left_value': ['a', 'b', 'c', 'd', 'e']})

right_frame = pd.DataFrame({'key': range(2, 7), 'right_value': ['f', 'g', 'h', 'i', 'j']})

left_frame
right_frame

# inner join
pd.merge(left_frame, right_frame, on='key', how='inner')

# left outer join
pd.merge(left_frame, right_frame, on='key', how='left')

# right outer join
pd.merge(left_frame, right_frame, on='key', how='right')

# full outer join
pd.merge(left_frame, right_frame, on='key', how='outer')
```

27

Combining Data Frames (concat)

- Similar to the SQL union clause.
- `pandas.concat` takes a list of Series or DataFrames and returns a Series or DataFrame of the concatenated objects.
- can specify many objects to combine simultaneously (however need at least 2!).

```
pd.concat([left_frame, right_frame])
```


- The default is that the objects are vertically appended.
- Columns with the same name will be combined.
- To combine objects side-by-side, this can be easily specified using the `axis` option, e.g.

```
pd.concat([left_frame, right_frame], axis=1)
```

28

Grouping Data Frames (groupby)

- Used to group data in some meaningful way, so that we can perform operations over each separate group (e.g. calculating the average expenditure on customer purchases, grouped by product type (clothing, books, DVDs, electronics, etc.)).
- This much-referenced graphic explains what's happening :)



Hadley Wickham
Data Science in R

- A very useful feature once you get your head around it!

29

Grouping Data Frames (groupby)

```
import pandas as pd

# Read in the data and check it out
mtcars = pd.read_csv("mtcars.csv")
mtcars.head()
mtcars.shape

# Compute basic descriptive stats over data
mtcars.describe()
mtcars.mean() # also compute median, std, var, min, max, quantile
mtcars.mean(axis=1) # compute row means (ie across columns)

# How many automatic transmission cars are there?
mtcars[mtcars["am"]==0].shape
mtcars[mtcars["am"]==0]

# Plot a histogram
mtcars["mpg"].hist()

# Group by number of carburettors and describe
grouped_by_carb = mtcars.groupby("carb")
grouped_by_carb.mean()
```

30

Grouping Data Frames (groupby)

```
# can group by more than one category
grouped_by_carb_am = mtcars.groupby(["carb", "am"])

# compute statistics aggregated over groupings
import numpy as np
grouped_by_carb_am.agg([np.mean, np.std])

# count the number of cars in each combination of carb and am
counts = grouped_by_carb_am['carb'].count()

# plot the counts
import matplotlib.pyplot as plt
%matplotlib inline
df = counts.unstack()
ax = df.plot(kind='bar', stacked=True, figsize=(20, 10),
             colormap="BuGn")
ax.set_ylabel("Count")
patches, labels = ax.get_legend_handles_labels()
ax.legend(patches, labels, loc='best')
```

31

Arithmetic between DataFrames and Series

- Can take advantage of **broadcasting** to perform operations between DataFrames and Series structures.

```
In [135]:
arr = np.arange(12.).reshape((3, 4))
```

```
Out[135]:
array([[ 0.,  1.,  2.,  3.],
       [ 4.,  5.,  6.,  7.],
       [ 8.,  9., 10., 11.]])
```

```
In [136]:
arr - arr[0]
```

```
Out[136]:
array([[ 0.,  0.,  0.,  0.],
       [ 4.,  4.,  4.,  4.],
       [ 8.,  8.,  8.,  8.]])
```

32

Arithmetic between DataFrames and Series

- By default, arithmetic between DataFrame and Series matches the index of the Series on the DataFrame's columns, **broadcasting down the rows**:

```
frame = pd.DataFrame(np.arange(12.).reshape((4, 3)),
                      columns=list('bde'), index=['Utah', 'Ohio', 'Texas', 'Oregon'])
series = frame.ix[0]
```

frame

| | b | d | e |
|--------|---|----|----|
| Utah | 0 | 1 | 2 |
| Ohio | 3 | 4 | 5 |
| Texas | 6 | 7 | 8 |
| Oregon | 9 | 10 | 11 |

series

```
b    0
d    1
e    2
Name: Utah, dtype: float64
```

| | b | d | e |
|--------|---|---|---|
| Utah | 0 | 0 | 0 |
| Ohio | 3 | 3 | 3 |
| Texas | 6 | 6 | 6 |
| Oregon | 9 | 9 | 9 |

```
In[138]: frame - series
```

33

Arithmetic between DataFrames and Series

- To **broadcast down the columns** and match on the rows, need to use an arithmetic method:

```
series3 = frame['d']
frame.sub(series3, axis=0)
```

frame

| | b | d | e |
|--------|---|----|----|
| Utah | 0 | 1 | 2 |
| Ohio | 3 | 4 | 5 |
| Texas | 6 | 7 | 8 |
| Oregon | 9 | 10 | 11 |

frame - series3

| | b | d | e |
|--------|----|---|---|
| Utah | -1 | 0 | 1 |
| Ohio | -1 | 0 | 1 |
| Texas | -1 | 0 | 1 |
| Oregon | -1 | 0 | 1 |

34

A Closer Look At Indexing

- Any array, or other sequence of labels used when constructing a Series or DataFrame, is internally converted to an Index:

```
In [50]: obj = pd.Series(range(3), index=['a', 'b', 'c'])
```

- Index objects are immutable and thus can't be modified by the user:

```
In [51]: obj.index[1] = 'd'
```

```
-----
NameError                                Traceback (most recent call last)
.
.
.
<class 'pandas.core.index.Index'> does not support mutable
operations.
```

- Immutability is important so that Index objects can be safely shared among data structures. However, we can kind of get around this ...

35

Reindexing

- A critical method on pandas objects is `reindex`, which means to create a new object with the data conformed to a new index.

```
In [56]:
obj = pd.Series([4.5, 7.2, -5.3, 3.6],
                index=['d', 'b', 'a', 'c'])
```

- Calling `reindex` on this Series rearranges the data according to the new index, introducing missing values if any index values were not already present:

```
In [57]:
obj2 = obj.reindex(['a', 'b', 'c', 'd', 'e'])
obj2
```

```
Out[57]:
a -5.3
b 7.2
c 3.6
d 4.5
e NaN
```

36

Reindexing

- Can avoid the usage of NaN when filling in an empty index by using `fill_value`:

```
In [61]:
obj = pd.Series([4.5, 7.2, -5.3, 3.6],
index=['d', 'b', 'a', 'c'])
obj
```

```
Out[61]:
d 4.5
b 7.2
a -5.3
c 3.6
```

- If we re-index, we can initialise all non-existing indices:

```
In [63]:
obj.reindex(['a', 'b', 'c', 'd', 'e'], fill_value=0)
```

```
Out[63]:
a -5.3
b 7.2
c 3.6
d 4.5
e 0.0
```

37

Reindexing

- Can also fill empty values associated with indices, using non-obvious (but useful) methods!

```
In [65]: obj3 = pd.Series(['blue', 'purple', 'yellow'],
index=[0, 2, 4])
```

```
Out[65]:
0 blue
2 purple
4 yellow
dtype: object
```

```
In [66]: obj3 = obj3.reindex(range(6), method='ffill')
```

```
Out[66]:
0 blue
1 blue
2 purple
3 purple
4 yellow
5 yellow
```

reindex options

| Argument | Description |
|---|---------------------------------|
| <code>ffill</code> or <code>pad</code> | Fill (or carry) values forward |
| <code>bfill</code> or <code>backfill</code> | Fill (or carry) values backward |

38

A Closer Look At Reindexing

- reindex can alter either the (row) index, columns, or both.

```
frame = pd.DataFrame(np.arange(9).reshape((3, 3)), index=['a', 'c', 'd'],  
columns=['Ohio', 'Texas', 'California'])
```

| | Ohio | Texas | California |
|---|------|-------|------------|
| a | 0 | 1 | 2 |
| c | 3 | 4 | 5 |
| d | 6 | 7 | 8 |

- Specifying just a sequence will re-order the rows (by default).

```
frame2 = frame.reindex(['a', 'b', 'c', 'd'])
```

| | Ohio | Texas | California |
|---|------|-------|------------|
| a | 0 | 1 | 2 |
| b | NaN | NaN | NaN |
| c | 3 | 4 | 5 |
| d | 6 | 7 | 8 |

A Closer Look At Reindexing

- The columns can be reindexed using the columns keyword:

```
states = ['Texas', 'Utah', 'California']  
frame.reindex(columns=states)
```

| | Texas | Utah | California |
|---|-------|------|------------|
| a | 1 | NaN | 2 |
| c | 4 | NaN | 5 |
| d | 7 | NaN | 8 |

- Can reindex both simultaneously:

```
frame.reindex(index=['a', 'b', 'c', 'd'],  
method='ffill', columns=states)
```

| | Texas | Utah | California |
|---|-------|------|------------|
| a | 1 | NaN | 2 |
| b | 1 | NaN | 2 |
| c | 4 | NaN | 5 |
| d | 7 | NaN | 8 |

A Closer Look At Reindexing

- reindexing can be done more succinctly by label-indexing with `ix`:

```
frame.ix[['a', 'b', 'c', 'd'], states]
```

reindex function arguments

| Argument | Description |
|-------------------------|---|
| <code>index</code> | New sequence to use as index. Can be Index instance or any other sequence-like Python data structure. An Index will be used exactly as is without any copying |
| <code>method</code> | Interpolation (fill) method |
| <code>fill_value</code> | Substitute value to use when introducing missing data by reindexing |
| <code>limit</code> | When forward- or backfilling, maximum size gap to fill |
| <code>level</code> | Match simple Index on level of MultiIndex, otherwise select subset of |
| <code>copy</code> | Do not copy underlying data if new index is equivalent to old index. True by default (i.e. always copy data). |

Indexing: a few things to watch out for!

- Series indexing works as NumPy indexing**, except can use the Series index values in place of just the integers.

```
In[94]:  
obj = pd.Series(np.arange(4.), index=['a', 'b', 'c', 'd'])  
print(obj[1], ' ', obj['b'])
```

```
Out[94]:  
1.0 1.0
```

- Can use slices, boolean indexing, etc., as seen with NumPy.
- However slicing with labels in pandas takes the endpoint (unlike base Python and NumPy). This matches R's functionality.

```
In[95]: obj['b':'c']  
Out[95]:  
b      1  
c      2
```

Indexing: a few things to watch out for!

- **DataFrame indexing** allows the retrieval of one or more columns either with a single value or sequence:

```
data =  
pd.DataFrame(np.arange(16).reshape((4, 4)),  
index=['Ohio', 'Colorado', 'Utah', 'New York'],  
columns=['one', 'two', 'three', 'four'])
```

```
data['two']
```

```
Out[103]:  
Ohio      1  
Colorado  5  
Utah      9  
New York 13
```

```
data[['two', 'four']]
```

| | two | four |
|----------|-----|------|
| Ohio | 1 | 3 |
| Colorado | 5 | 7 |
| Utah | 9 | 11 |
| New York | 13 | 15 |

- However there are also a few special cases ...

43

Indexing: a few things to watch out for!

- **Slicing with labels behaves as with Series slicing.** Slicing with row numbers however:

```
In[120]: data[:2]
```

| | one | two | three | four |
|----------|-----|-----|-------|------|
| Ohio | 0 | 1 | 2 | 3 |
| Colorado | 4 | 5 | 6 | 7 |

```
In[121]: data[data['three'] > 5]
```

| | one | two | three | four |
|----------|-----|-----|-------|------|
| Colorado | 4 | 5 | 6 | 7 |
| Utah | 8 | 9 | 10 | 11 |
| New York | 12 | 13 | 14 | 15 |

- *Slicing with indices behaves as in base Python!*

- Can also use a pandas structure in a boolean statement:

```
In[126]: data < 5  
In[127]: data[data<5]=0 # Try it!
```

| | one | two | three | four |
|----------|-------|-------|-------|-------|
| Ohio | True | True | True | True |
| Colorado | True | False | False | False |
| Utah | False | False | False | False |
| New York | False | False | False | False |

44

DataFrame and ix

- Label-indexing of DataFrames can also use the `ix` function

```
In[127]: data.ix['Colorado', ['two', 'three']]
```

```
In [128]: data.ix[['Colorado', 'Utah'], ['two', 'three']]
```

```
In[129]: data.ix[2]
```

```
In[130]: data.ix[:, 'Utah', 'two']
```

```
In[131]: data.ix[data.three > 5, :3]
```

Note use of `.` operator to access column variable – very R-like!

45

Hierarchical Indexing: Series

- Enables multiple (two or more) index levels on an axis.
- Provides a way to work with higher dimensional data at lower dimensions. Mimics R functionality.

```
data = pd.Series(np.random.randn(10),  
index=[['a', 'a', 'a', 'b', 'b', 'b', 'c', 'c', 'd', 'd'],  
       [1, 2, 3, 1, 2, 3, 1, 2, 2, 3]])
```

```
Out[167]: a 1    1.612344  
          2   -1.212449  
          3    2.552081  
         b 1   -2.021124  
          2   -0.156685  
          3   -0.458148  
         c 1   -1.592207  
          2    0.144602  
         d 2   -1.991424  
          3   -0.110237
```

46

Hierarchical Indexing: Series

- Makes **partial indexing** possible.

```
data['b']
```

```
Out[169]: 1    -2.021124
          2    -0.156685
          3    -0.458148
          dtype: float64
```

```
data['b':'c'] # or data.ix[['b', 'c']]
```

```
Out[170]: b 1    -2.021124
          2    -0.156685
          3    -0.458148
          c 1    -1.592207
          2     0.144602
          dtype: float64
```

```
data[:,2]
```

```
Out[172]: a    -1.212449
          b    -0.156685
          c     0.144602
          d    -1.991424
          dtype: float64
```

47

Hierarchical Indexing: Series stack and unstack

- Another replication of R functionality.

```
Out[173]:
```

```
data.stack() #notice NaNs
```

| | 1 | 2 | 3 |
|---|-----------|-----------|-----------|
| a | 1.612344 | -1.212449 | 2.552081 |
| b | -2.021124 | -0.156685 | -0.458148 |
| c | -1.592207 | 0.144602 | NaN |
| d | NaN | -1.991424 | -0.110237 |

```
data.stack().unstack()
```

```
Out[167]: a 1     1.612344
          2    -1.212449
          3     2.552081
          b 1    -2.021124
          2    -0.156685
          3    -0.458148
          c 1    -1.592207
          2     0.144602
          d 2    -1.991424
          3    -0.110237
```


Hierarchical Indexing: DataFrames

- With DataFrames, either axis can have a hierarchical index.

```
frame = pd.DataFrame(np.arange(12).reshape((4, 3)), index=[['a',  
'a', 'b', 'b'], [1, 2, 1, 2]], columns=[['Ohio', 'Ohio',  
'Colorado'], ['Green', 'Red', 'Green']])
```

| | | Ohio | | Colorado |
|---|---|-------|-----|----------|
| | | Green | Red | Green |
| a | 1 | 0 | 1 | 2 |
| | 2 | 3 | 4 | 5 |
| b | 1 | 6 | 7 | 8 |
| | 2 | 9 | 10 | 11 |

49

Hierarchical Indexing: DataFrames

- With DataFrames, either axis can have a hierarchical index.

```
frame.index.names = ['key1', 'key2']  
frame.columns.names = ['state', 'colour']
```

| | state | Ohio | | Colorado |
|------|--------|-------|-----|----------|
| | colour | Green | Red | Green |
| key1 | key2 | | | |
| a | 1 | 0 | 1 | 2 |
| | 2 | 3 | 4 | 5 |
| b | 1 | 6 | 7 | 8 |
| | 2 | 9 | 10 | 11 |

| | colour | Green | Red |
|------|--------|-------|-----|
| key1 | key2 | | |
| a | 1 | 0 | 1 |
| | 2 | 3 | 4 |
| b | 1 | 6 | 7 |
| | 2 | 9 | 10 |

```
frame['Ohio'] # partial indexing
```

50

DataFrame Indexing Options

| Type | Notes |
|---|--|
| <code>obj[val]</code> | Select single column or sequence of columns from the DataFrame. Special case conveniences: boolean array (filter rows), slice (slice rows), or boolean DataFrame (set values based on some criterion). |
| <code>obj.ix[val]</code> | Selects single row of subset of rows from the DataFrame. |
| <code>obj.ix[:, val]</code> | Selects single column of subset of columns. |
| <code>obj.ix[val1, val2]</code> | Select both rows and columns. |
| <code>reindex</code> method | Conform one or more axes to new indexes. |
| <code>xs</code> method | Select single row or column as a Series by label. |
| <code>icol</code> , <code>irow</code> methods | Select single column or row, respectively, as a Series by integer location. |
| <code>get_value</code> , <code>set_value</code> methods | Select single value by row and column label. |

51

Function mapping: apply

- In R, there is much usage of a suite of apply functions (apply, sapply, lapply...). We'll see these in the next lecture!
- Python-defined ufuncs 'automatically' apply themselves to each element in a pandas DataFrame. However user-defined functions do not.
- pandas replicates the R functionality in part by providing an apply function for this purpose.

52

Function mapping: apply

```
frame = pd.DataFrame(np.random.randn(4, 3), columns=list('bde'),  
index=['Utah', 'Ohio', 'Texas', 'Oregon'])
```

| | b | d | e |
|--------|-----------|-----------|-----------|
| Utah | -1.284567 | -0.724012 | -0.410331 |
| Ohio | -1.028510 | 0.117753 | -0.048079 |
| Texas | -1.210197 | -2.220176 | -1.135762 |
| Oregon | 0.479340 | -0.005978 | -1.975791 |

```
f = lambda x: x.max() - x.min()  
frame.apply(f)  
b 2.142073  
d 2.710013  
e 0.349985
```

```
frame.apply(f,axis=1)  
Utah 2.550449  
Ohio 1.386588  
Texas 0.905827  
Oregon 1.248064
```

Many of the most common array statistics (e.g. sum, mean) are DataFrame methods, so using apply is not necessary. (in R this is not the case!)

53

Descriptive Statistics (pandas)

| Method | Description |
|----------------|---|
| count | Number of non-NA values |
| describe | Compute set of summary statistics for Series or each DataFrame column |
| min, max | Compute minimum and maximum values |
| argmin, argmax | Compute index locations (integers) at which minimum or maximum value obtained, respectively |
| idxmin, idxmax | Compute index values at which minimum or maximum value obtained, respectively |
| quantile | Compute sample quantile ranging from 0 to 1 |
| sum | Sum of values |
| mean | Mean of values |
| median | Arithmetic median (50% quantile) of values |
| mad | Mean absolute deviation from mean value |
| var | Sample variance of values |
| std | Sample standard deviation of values |
| skew | Sample skewness (3rd moment) of values |
| kurt | Sample kurtosis (4th moment) of values |

4

Descriptive Statistics (pandas)

| | |
|----------------|--|
| cumsum | Cumulative sum of values |
| cummin, cummax | Cumulative minimum or maximum of values, respectively |
| cumprod | Cumulative product of values |
| diff | Compute 1st arithmetic difference (useful for time series) |
| pct_change | Compute percent changes |

55

apply and applymap

```
def f(x): # returns a series
    return Series([x.min(), x.max()], index=['min', 'max'])
```

```
frame.apply(f)
```

| | b | d | e |
|-----|-----------|----------|-----------|
| min | -0.555730 | 0.281746 | -1.296221 |
| max | 1.246435 | 1.965781 | 1.393406 |

```
format = lambda x: # returns a formatted string
    '%.2f' % x
```

```
frame.applymap(format)
```

| | b | d | e |
|--------|-------|------|-------|
| Utah | -0.20 | 0.48 | -0.52 |
| Ohio | -0.56 | 1.97 | 1.39 |
| Texas | 0.09 | 0.28 | 0.77 |
| Oregon | 1.25 | 1.01 | -1.30 |

applymap is used for element-wise function application

56

Missing Data

- We can fill in, or filter out missing data (NAs). The functionalities available are:

| Argument | Description |
|----------------------|---|
| <code>dropna</code> | Filter axis labels based on whether values for each label have missing data, with varying thresholds for how much missing data to tolerate. |
| <code>fillna</code> | Fill in missing data with some value or using an interpolation method such as 'ffill' or 'bfill'. |
| <code>isnull</code> | Return like-type object containing boolean values indicating which values are missing / NA. |
| <code>notnull</code> | Negation of <code>isnull</code> . |

```
from numpy import nan as NA
data = Series([1, NA, 3.5, NA, 7])
data.dropna()    # data[data.notnull()] gives the same result

0 1
2 3.5
4 7
```

57

Detecting Missing Data in Series

- Can specify indices when passing a dict structure in `Series()`:

```
In [22]: states = ['California', 'Ohio', 'Oregon', 'Texas']
        obj4 = pd.Series(sdata, index=states)
```

```
Out[22]:
California NaN
Ohio 35000
Oregon 16000
Texas 71000
dtype: float64
```

*Introduces a missing number NaN
(indicates NA values)*

- Use `isnull` and `notnull` to detect missing data:

```
In [26]: pd.isnull(obj4)
Out[26]:
California    True
Ohio         False
Oregon       False
Texas        False
```

```
In [27]: pd.notnull(obj4)
Out[27]:
California    False
Ohio         True
Oregon       True
Texas        True
```

58

Filling In Missing Data in Series

- We can use `fillna()` to do the job.
- Recall:

```
states = ['California', 'Ohio', 'Oregon', 'Texas']  
obj4 = pd.Series(sdata, index=states)
```

- Use `fillna()` to replace the NaNs.

```
In [8]:obj4.fillna(0)
```

```
Out[8]: California      0  
        Ohio           35000  
        Oregon         16000  
        Texas          71000
```

- **Caution!** Simply replacing NaNs because they are inconvenient is very poor practice - you should never alter data for ease. You should only fill in missing values when you know why they occur and how to replace them.
- Otherwise we are better off ignoring the rows of data containing NAs 59 and focusing on data we can trust.

Filling In Missing Data

- Can use `fillna()` in a number of ways:

```
df.fillna(0) # fill with 0s
```

```
df.fillna({1: 0.5, 3: -1}) # fill with a dict.  
# Non-existent indices ignored
```

```
_ = df.fillna(0, inplace=True) # fill in place (no copy created)
```

```
df.fillna(method='bfill') # can use the reindexing methods
```

```
data = Series([1., NA, 3.5, NA, 7]) # fill with the mean  
data.fillna(data.mean())
```

Filling In Missing Data

| Argument | Description |
|----------|---|
| value | Scalar value or dict-like object to use to fill missing values |
| method | Interpolation, by default 'ffill' if function called with no other arguments |
| axis | Axis to fill on, default axis=0 |
| inplace | Modify the calling object without producing a copy |
| limit | For forward and backward filling, maximum number of consecutive periods to fill |

61

DataFrame – Inserting Missing Data

- To avoid issues with arithmetic operations, may wish to fill NaNs with a special value (e.g. 0) in these cases.

```
df1 = pd.DataFrame(np.arange(12.).reshape((3, 4)), columns=list('abcd'))  
df2 = pd.DataFrame(np.arange(20.).reshape((4, 5)), columns=list('abcde'))
```

```
df1.add(df2, fill_value=0)
```

Out[134]:

| | a | b | c | d | e |
|---|----|----|----|----|----|
| 0 | 0 | 2 | 4 | 6 | 4 |
| 1 | 9 | 11 | 13 | 15 | 9 |
| 2 | 18 | 20 | 22 | 24 | 14 |
| 3 | 15 | 16 | 17 | 18 | 19 |

- Can use `fill_value` with `add`, `sub`, `div` and `mul`.

62

Filtering Missing Data

- DataFrames are a bit more complex.

```
data = DataFrame([[1., 6.5, 3.], [1., NA, NA], [NA, NA, NA],  
[NA, 6.5, 3.]])
```

```
data
```

| | 0 | 1 | 2 |
|---|-----|-----|-----|
| 0 | 1. | 6.5 | 3. |
| 1 | 1 | NaN | NaN |
| 2 | NaN | NaN | NaN |
| 3 | NaN | 6.5 | 3 |

```
cleaned = data.dropna()
```

| | 0 | 1 | 2 |
|---|----|-----|----|
| 0 | 1. | 6.5 | 3. |

```
cleaned = data.dropna(how='all') # returns NaN-only rows
```

```
cleaned = data.dropna(axis=1,how='all') # drops columns
```

63

Filtering Missing Data

- Can also use a threshold argument `thresh`

```
df = pd.DataFrame(np.random.randn(7, 3))  
df.ix[:4, 1] = NA; df.ix[:2, 2] = NA  
df
```

| | 0 | 1 | 2 |
|---|-----------|-----------|-----------|
| 0 | -0.148317 | NaN | NaN |
| 1 | -0.359840 | NaN | NaN |
| 2 | 0.891757 | NaN | NaN |
| 3 | 0.944898 | NaN | -0.064449 |
| 4 | -1.971901 | NaN | -1.366681 |
| 5 | 0.731808 | -0.807611 | -1.414878 |
| 6 | -1.406918 | 0.265153 | 1.455532 |

```
df.dropna(thresh=2)
```

| | 0 | 1 | 2 |
|---|-----------|-----------|-----------|
| 3 | 0.944898 | NaN | -0.064449 |
| 4 | -1.971901 | NaN | -1.366681 |
| 5 | 0.731808 | -0.807611 | -1.414878 |
| 6 | -1.406918 | 0.265153 | 1.455532 |

64